

Partitioned tables in BigQuery

Introduction to partitioned tables

1. A partitioned table is a special table that is divided into segments, called partitions, that make it easier to manage and query your data.
2. By dividing a large table into smaller partitions, you can improve query performance, and you can control costs by reducing the number of bytes read by a query.

You can partition BigQuery tables by:

- **Time-unit column:** Tables are partitioned based on a [TIMESTAMP](#), [DATE](#), or [DATETIME](#) column in the table. BigQuery automatically puts the data into the correct partition, based on the values in the column.

For [TIMESTAMP](#) and [DATETIME](#) columns, the partitions can have either hourly, daily, monthly, or yearly granularity. For [DATE](#) columns, the partitions can have daily, monthly, or yearly granularity. Partitions boundaries are based on UTC time

- **Ingestion time:** Tables are partitioned based on the timestamp when BigQuery ingests the data.
- **Integer range:** Tables are partitioned based on an integer column.

Creating partitioned tables

Create an empty partitioned table

The steps to create a partitioned table in BigQuery are similar to creating a [standard table](#), except that you specify the partitioning options, along with any other table options.

1. Create a time-unit column-partitioned table

To create an empty time-unit column-partitioned table with a schema definition:

1. Open the BigQuery page in the Cloud Console.
[Go to the BigQuery page](#)
2. In the **Explorer** panel, expand your project and select a dataset.
3. Expand the more_vert **Actions** option and click **Open**.
4. In the details panel, click **Create table** add_box.
5. On the **Create table** page, in the **Source** section, select **Empty table**.
6. In the **Destination** section:
 - For **Dataset name**, choose the appropriate dataset.
 - In the **Table name** field, enter the name of the table.
 - Verify that **Table type** is set to **Native table**.
7. In the **Schema** section, enter the schema definition. Make sure the schema includes a DATE, TIMESTAMP, or DATETIME column for the partitioning column. For more information, see [Specifying a schema](#).
8. In the **Partition and cluster settings** section, in the **Partitioning** drop-down list, select **Partition by field** and choose the partitioning column. This option is only available if the schema contains a DATE, TIMESTAMP, or DATETIME column.
9. Select the **Partitioning type** to choose daily, hourly, monthly, or yearly partitioning.
10. (Optional) To require a partition filter on all queries for this table, select the **Require partition filter** checkbox. Requiring a partition filter can reduce cost and improve performance. For more information, see [Set partition filter requirements](#).
11. Click **Create table**.

OR

Create table using SQL

```
CREATE TABLE
  mydataset.newtable (transaction_id INT64, transaction_date DATE)
PARTITION BY
  transaction_date
OPTIONS(
  partition_expiration_days=3,
  require_partition_filter=true
)
```

2. Create an ingestion-time partitioned table

To create an empty ingestion-time partitioned table with a schema definition:

1. Open the BigQuery page in the Cloud Console.
[Go to the BigQuery page](#)
2. In the **Explorer** panel, expand your project and select a dataset.
3. Expand the more_vert **Actions** option and click **Open**.
4. In the details panel, click **Create table** add_box.
5. On the **Create table** page, in the **Source** section, select **Empty table**.
6. In the **Destination** section:
 - For **Dataset name**, choose the appropriate dataset.
 - In the **Table name** field, enter the name of the table.
 - Verify that **Table type** is set to **Native table**.
7. In the **Schema** section, enter the [schema](#) definition.
8. In the **Partition and cluster settings** section, for **Partitioning**, click **Partition by ingestion time**.
9. (Optional) To require a partition filter on all queries for this table, select the **Require partition filter** checkbox. Requiring a partition filter can reduce cost and improve performance. For more information, see [Set partition filter requirements](#).
10. Click **Create table**.

3. Create an integer-range partitioned table

To create an empty integer-range partitioned table with a schema definition:

1. Open the BigQuery page in the Cloud Console.
[Go to the BigQuery page](#)
2. In the **Explorer** panel, expand your project and select a dataset.
3. Expand the more_vert **Actions** option and click **Open**.
4. In the details panel, click **Create table** add_box.
5. On the **Create table** page, in the **Source** section, select **Empty table**.
6. In the **Destination** section:
 - For **Dataset name**, choose the appropriate dataset.
 - In the **Table name** field, enter the name of the table.
 - Verify that **Table type** is set to **Native table**.
7. In the **Schema** section, enter the schema definition. Make sure the schema includes an INTEGER column for the partitioning column. For more information, see [Specifying a schema](#).
8. In the **Partition and cluster settings** section, in the **Partitioning** drop-down list, select **Partition by field** and choose the partitioning column. This option is only available if the schema contains an INTEGER column.
9. Provide values for **Start**, **End**, and **Interval**:
 - **Start** is the start of first partition range (inclusive).
 - **End** is the end of last partition range (exclusive).
 - **Interval** is the width of each partition range.
10. Values outside of these ranges go into a special __UNPARTITIONED__ partition.
11. (Optional) To require a partition filter on all queries for this table, select the **Require partition filter** checkbox. Requiring a partition filter can reduce cost and improve performance. For more information, see [Set partition filter requirements](#).
12. Click **Create table**.

Managing partitioned tables

Getting partition metadata using INFORMATION_SCHEMA views

When you query the INFORMATION_SCHEMA.PARTITIONS view, the query results contain one row for each partition. For example, the following query lists all of the table partitions in the dataset named mydataset:

```
SELECT table_name, partition_id, total_rows  
FROM `mydataset.INFORMATION_SCHEMA.PARTITIONS`  
WHERE partition_id IS NOT NULL
```