

Joining Data Frames

Let us understand how to join Data Frames using Pandas.

```
%run 06_csv_to_pandas_data_frame.ipynb
```

orders

	order_id	order_date	order_customer_id	order_status
0	1	2013-07-25 00:00:00.0	11599	CLOSED
1	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT
2	3	2013-07-25 00:00:00.0	12111	COMPLETE
3	4	2013-07-25 00:00:00.0	8827	CLOSED
4	5	2013-07-25 00:00:00.0	11318	COMPLETE
...
68878	68879	2014-07-09 00:00:00.0	778	COMPLETE
68879	68880	2014-07-13 00:00:00.0	1117	COMPLETE
68880	68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT
68881	68882	2014-07-22 00:00:00.0	10000	ON_HOLD
68882	68883	2014-07-23 00:00:00.0	5533	COMPLETE

68883 rows × 4 columns

order_items

	order_item_id	order_item_order_id	order_item_product_id	order_item_quantity	order
0	1	1	957	1	
1	2	2	1073	1	
2	3	2	502	5	
3	4	2	403	1	
4	5	4	897	2	
...
172193	172194	68881	403	1	
172194	172195	68882	365	1	
172195	172196	68882	502	1	
172196	172197	68883	208	1	
172197	172198	68883	502	3	

172198 rows × 6 columns

- Join orders and order_items using orders.order_id and order_items.order_item_order_id.

```
orders.join?
```

Contents

[Task 1](#)

[Task 2](#)

[Task 3](#)

Print to PDF ▶

Signature: `orders.join(other, on=None, how='left', lsuffix='', rsuffix='', sort=False) -> 'DataFrame'`

Docstring:

Join columns of another DataFrame.

Join columns with `other` DataFrame either on index or on a key column. Efficiently join multiple DataFrame objects by index at once by passing a list.

Parameters

`other` : DataFrame, Series, or list of DataFrame

Index should be similar to one of the columns in this one. If a Series is passed, its name attribute must be set, and that will be used as the column name in the resulting joined DataFrame.

`on` : str, list of str, or array-like, optional

Column or index level name(s) in the caller to join on the index in `other`, otherwise joins index-on-index. If multiple values given, the `other` DataFrame must have a MultiIndex. Can pass an array as the join key if it is not already contained in the calling DataFrame. Like an Excel VLOOKUP operation.

`how` : {'left', 'right', 'outer', 'inner'}, default 'left'

How to handle the operation of the two objects.

* left: use calling frame's index (or column if on is specified)

* right: use `other`'s index.

* outer: form union of calling frame's index (or column if on is specified) with `other`'s index, and sort it. lexicographically.

* inner: form intersection of calling frame's index (or column if on is specified) with `other`'s index, preserving the order of the calling's one.

`lsuffix` : str, default ''

Suffix to use from left frame's overlapping columns.

`rsuffix` : str, default ''

Suffix to use from right frame's overlapping columns.

`sort` : bool, default False

Order result DataFrame lexicographically by the join key. If False, the order of the join key depends on the join type (how keyword).

Returns

DataFrame

A dataframe containing columns from both the caller and `other`.

See Also

`DataFrame.merge` : For column(s)-on-columns(s) operations.

Notes

Parameters `on`, `lsuffix`, and `rsuffix` are not supported when passing a list of `DataFrame` objects.

Support for specifying index levels as the `on` parameter was added in version 0.23.0.

Examples

```
>>> df = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K3', 'K4', 'K5'],
...                    'A': ['A0', 'A1', 'A2', 'A3', 'A4', 'A5']})
```

```
>>> df
   key  A
0  K0  A0
1  K1  A1
2  K2  A2
3  K3  A3
4  K4  A4
5  K5  A5
```

```
>>> other = pd.DataFrame({'key': ['K0', 'K1', 'K2'],
...                       'B': ['B0', 'B1', 'B2']})
```

```
>>> other
   key  B
0  K0  B0
1  K1  B1
2  K2  B2
```

Join DataFrames using their indexes.

```
>>> df.join(other, lsuffix='_caller', rsuffix='_other')
```

```
   key_caller  A key_other  B
0         K0  A0         K0  B0
1         K1  A1         K1  B1
2         K2  A2         K2  B2
3         K3  A3         NaN  NaN
```

4	K4	A4	NaN	NaN
5	K5	A5	NaN	NaN

If we want to join using the key columns, we need to set key to be the index in both `df` and `other`. The joined DataFrame will have key as its index.

```
>>> df.set_index('key').join(other.set_index('key'))
      A      B
key
K0  A0  B0
K1  A1  B1
K2  A2  B2
K3  A3  NaN
K4  A4  NaN
K5  A5  NaN
```

Another option to join using the key columns is to use the `on` parameter. DataFrame.join always uses `other`'s index but we can use any column in `df`. This method preserves the original DataFrame's index in the result.

```
>>> df.join(other.set_index('key'), on='key')
      key  A      B
0  K0  A0  B0
1  K1  A1  B1
2  K2  A2  B2
3  K3  A3  NaN
4  K4  A4  NaN
5  K5  A5  NaN
File:      /opt/anaconda3/envs/beakerx/lib/python3.6/site-packages/pandas/core/frame.py
Type:      method
```

```
orders.set_index('order_id')
```

	order_date	order_customer_id	order_status
order_id			
1	2013-07-25 00:00:00.0	11599	CLOSED
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT
3	2013-07-25 00:00:00.0	12111	COMPLETE
4	2013-07-25 00:00:00.0	8827	CLOSED
5	2013-07-25 00:00:00.0	11318	COMPLETE
...
68879	2014-07-09 00:00:00.0	778	COMPLETE
68880	2014-07-13 00:00:00.0	1117	COMPLETE
68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT
68882	2014-07-22 00:00:00.0	10000	ON_HOLD
68883	2014-07-23 00:00:00.0	5533	COMPLETE

68883 rows × 3 columns

```
order_items.set_index('order_item_order_id')
```

	order_item_id	order_item_product_id	order_item_quantity	order_item_sub
order_item_order_id				
1	1	957	1	29
2	2	1073	1	19
2	3	502	5	25
2	4	403	1	12
4	5	897	2	4
...	
68881	172194	403	1	12
68882	172195	365	1	5
68882	172196	502	1	5
68883	172197	208	1	199
68883	172198	502	3	15

172198 rows × 5 columns

```
# Join orders and order_items using order_id (order_item_order_id from order_items)
orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'))
```

	order_date	order_customer_id	order_status	order_item_id	order_item_product
1	2013-07-25 00:00:00.0	11599	CLOSED	1.0	95
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	2.0	107
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	3.0	50
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	4.0	40
3	2013-07-25 00:00:00.0	12111	COMPLETE	NaN	1
...	
68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT	172194.0	40
68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172195.0	36
68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172196.0	50
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172197.0	20
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172198.0	50

183650 rows × 8 columns

```
orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'), how='inner')
```

	order_date	order_customer_id	order_status	order_item_id	order_item_product_id
1	2013-07-25 00:00:00.0	11599	CLOSED	1	9
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	2	10
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	3	11
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	4	12
4	2013-07-25 00:00:00.0	8827	CLOSED	5	13
...
68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT	172194	14
68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172195	15
68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172196	16
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172197	17
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172198	18

172198 rows × 8 columns

```
orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'), how='inner'). \
    reset_index()
```

	index	order_date	order_customer_id	order_status	order_item_id	order_ite
0	1	2013-07-25 00:00:00.0	11599	CLOSED	1	
1	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	2	
2	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	3	
3	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	4	
4	4	2013-07-25 00:00:00.0	8827	CLOSED	5	
...	
172193	68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT	172194	
172194	68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172195	
172195	68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172196	
172196	68883	2014-07-23 00:00:00.0	5533	COMPLETE	172197	
172197	68883	2014-07-23 00:00:00.0	5533	COMPLETE	172198	

172198 rows × 9 columns

```
orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'), how='inner'). \
    reset_index(). \
    rename(columns={'index': 'order_id'})
```

	order_id	order_date	order_customer_id	order_status	order_item_id	order_i
0	1	2013-07-25 00:00:00.0	11599	CLOSED	1	
1	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	2	
2	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	3	
3	2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	4	
4	4	2013-07-25 00:00:00.0	8827	CLOSED	5	
...
172193	68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT	172194	
172194	68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172195	
172195	68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172196	
172196	68883	2014-07-23 00:00:00.0	5533	COMPLETE	172197	
172197	68883	2014-07-23 00:00:00.0	5533	COMPLETE	172198	

172198 rows × 9 columns

Task 1

Compute Daily Revenue using orders.order_date and order_items.order_item_order_subtotal considering only COMPLETE and CLOSED orders.

- Here are the steps to join orders and order_items and get daily revenue.
 - Create Data Frames for both orders and order_items using data in files.
 - Filter for orders which are either in **COMPLETE** or **CLOSED** status.
 - Set the join index for both the Data Frames.
 - Join both the Data Frames using **inner**.
 - Group the join results using **order_date** and get daily revenue by using **sum** on top of **order_item_subtotal**.

```
orders_considered = orders.query("order_status in ('COMPLETE', 'CLOSED')")
```

```
orders_filtered = orders[orders.order_status.isin(["COMPLETE", "CLOSED"])]
```

```
orders_considered. \
  set_index('order_id'). \
  join(order_items.set_index('order_item_order_id'), how='inner'). \
  groupby('order_date')['order_item_subtotal']. \
  agg(['sum']). \
  rename(columns={'sum': 'revenue'})
```

	revenue
order_date	
2013-07-25 00:00:00.0	31547.23
2013-07-26 00:00:00.0	54713.23
2013-07-27 00:00:00.0	48411.48
2013-07-28 00:00:00.0	35672.03
2013-07-29 00:00:00.0	54579.70
...	...
2014-07-20 00:00:00.0	60047.45
2014-07-21 00:00:00.0	51427.70
2014-07-22 00:00:00.0	36717.24
2014-07-23 00:00:00.0	38795.23
2014-07-24 00:00:00.0	50885.19

364 rows × 1 columns

Task 2

Get all the orders for which there are no corresponding order items.

- We can use default join (**left**) to get orders with out corresponding order items.

```
orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'))
```

	order_date	order_customer_id	order_status	order_item_id	order_item_produc
1	2013-07-25 00:00:00.0	11599	CLOSED	1.0	95
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	2.0	107
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	3.0	50
2	2013-07-25 00:00:00.0	256	PENDING_PAYMENT	4.0	40
3	2013-07-25 00:00:00.0	12111	COMPLETE	NaN	1
...
68881	2014-07-19 00:00:00.0	2518	PENDING_PAYMENT	172194.0	40
68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172195.0	36
68882	2014-07-22 00:00:00.0	10000	ON_HOLD	172196.0	50
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172197.0	20
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172198.0	50

183650 rows × 8 columns


```
orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id')). \
    query('order_item_id.isna()')
```

	order_date	order_customer_id	order_status	order_item_id	order_item_product_id	orc
3	2013-07-25 00:00:00.0	12111	COMPLETE	NaN	NaN	
6	2013-07-25 00:00:00.0	7130	COMPLETE	NaN	NaN	
22	2013-07-25 00:00:00.0	333	COMPLETE	NaN	NaN	
26	2013-07-25 00:00:00.0	7562	COMPLETE	NaN	NaN	
32	2013-07-25 00:00:00.0	3960	COMPLETE	NaN	NaN	
...	
68867	2014-06-23 00:00:00.0	869	CANCELED	NaN	NaN	
68872	2014-06-29 00:00:00.0	3354	COMPLETE	NaN	NaN	
68874	2014-07-03 00:00:00.0	1601	COMPLETE	NaN	NaN	
68876	2014-07-06 00:00:00.0	4124	COMPLETE	NaN	NaN	
68877	2014-07-07 00:00:00.0	9692	ON_HOLD	NaN	NaN	

11452 rows × 8 columns

```
orders_joined = orders.set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'))
```

```
orders_joined[orders_joined['order_item_id'].isna()]
```

	order_date	order_customer_id	order_status	order_item_id	order_item_product_id	orc
3	2013-07-25 00:00:00.0	12111	COMPLETE	NaN	NaN	
6	2013-07-25 00:00:00.0	7130	COMPLETE	NaN	NaN	
22	2013-07-25 00:00:00.0	333	COMPLETE	NaN	NaN	
26	2013-07-25 00:00:00.0	7562	COMPLETE	NaN	NaN	
32	2013-07-25 00:00:00.0	3960	COMPLETE	NaN	NaN	
...	
68867	2014-06-23 00:00:00.0	869	CANCELED	NaN	NaN	
68872	2014-06-29 00:00:00.0	3354	COMPLETE	NaN	NaN	
68874	2014-07-03 00:00:00.0	1601	COMPLETE	NaN	NaN	
68876	2014-07-06 00:00:00.0	4124	COMPLETE	NaN	NaN	
68877	2014-07-07 00:00:00.0	9692	ON_HOLD	NaN	NaN	

11452 rows × 8 columns

Task 3

Compute Daily Product Revenue using orders.order_date as well as order_items.order_item_product_id and order_items.order_item_order_subtotal considering only COMPLETE and CLOSED orders.

```
orders_considered = orders.query("order_status in ('COMPLETE', 'CLOSED')")
```

```
orders_filtered = orders[orders.order_status.isin(["COMPLETE", "CLOSED"])]
```

```
orders_considered. \
    set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'), how='inner')
```

	order_date	order_customer_id	order_status	order_item_id	order_item_product_id	order_item_subtotal
1	2013-07-25 00:00:00.0	11599	CLOSED	1	957	
4	2013-07-25 00:00:00.0	8827	CLOSED	5	897	
4	2013-07-25 00:00:00.0	8827	CLOSED	6	365	
4	2013-07-25 00:00:00.0	8827	CLOSED	7	502	
4	2013-07-25 00:00:00.0	8827	CLOSED	8	1014	
...
68880	2014-07-13 00:00:00.0	1117	COMPLETE	172191	1073	
68880	2014-07-13 00:00:00.0	1117	COMPLETE	172192	1014	
68880	2014-07-13 00:00:00.0	1117	COMPLETE	172193	1014	
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172197	208	
68883	2014-07-23 00:00:00.0	5533	COMPLETE	172198	502	

75408 rows × 8 columns

```
orders_considered. \
  set_index('order_id'). \
  join(order_items.set_index('order_item_order_id'), how='inner'). \
  groupby(['order_date', 'order_item_product_id'])['order_item_subtotal']
```

<pandas.core.groupby.generic.SeriesGroupBy object at 0x7f6dc89feeb8>

```
list(orders_considered. \
  set_index('order_id'). \
  join(order_items.set_index('order_item_order_id'), how='inner'). \
  groupby(['order_date', 'order_item_product_id'])['order_item_subtotal'][:10])
```

```
[(('2013-07-25 00:00:00.0', 24),
  57762      319.96
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 93),
  17       74.97
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 134),
  12      100.0
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 191),
  12      499.95
  28      399.96
  28       99.99
  61      399.96
  71      499.95
  101     99.99
  57757   499.95
  57764   499.95
  57768   499.95
  57776    99.99
  57776    99.99
  57779   499.95
  57782   199.98
  57788   199.98
  57788   499.95
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 226),
  68691    599.99
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 365),
  4       299.95
  5       299.95
  15      179.97
  17      239.96
  18      119.98
  28       59.99
  37       59.99
  45       59.99
  57      179.97
  57      119.98
  61      119.98
  61      119.98
  71      119.98
  91      299.95
  57756    59.99
  57757   119.98
  57779   299.95
  57781   299.95
  57788   239.96
  67416    59.99
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 403),
  5       129.99
  18       129.99
  24       129.99
  35       129.99
  57       129.99
  88       129.99
  98       129.99
  57754   129.99
  57756   129.99
  57757   129.99
  57762   129.99
  57762   129.99
  57768   129.99
  57788   129.99
  68691   129.99
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 502),
  4       150.0
  12      250.0
  15       50.0
  24       50.0
  24      250.0
  51       50.0
  62       50.0
  67      150.0
  98      100.0
  57757   150.0
  57758    50.0
  57758   100.0
  57764   150.0
  67416   100.0
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 572),
  72      119.97
  Name: order_item_subtotal, dtype: float64),
 (('2013-07-25 00:00:00.0', 625),
```

```
57764    199.99
Name: order_item_subtotal, dtype: float64]
```

```
orders_considered. \
    set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'), how='inner'). \
    groupby(['order_date', 'order_item_product_id'])['order_item_subtotal']. \
    agg(['sum']). \
    rename(columns={'sum': 'revenue'})
```

		revenue
order_date	order_item_product_id	
2013-07-25 00:00:00.0	24	319.96
	93	74.97
	134	100.00
	191	5099.49
	226	599.99
...
2014-07-24 00:00:00.0	926	31.98
	957	5399.64
	1004	10399.48
	1014	3148.74
	1073	4199.79

9120 rows × 1 columns

```
orders_considered. \
    set_index('order_id'). \
    join(order_items.set_index('order_item_order_id'), how='inner'). \
    groupby(['order_date', 'order_item_product_id'])['order_item_subtotal']. \
    agg(['sum']). \
    rename(columns={'sum': 'revenue'}). \
    reset_index()
```

	order_date	order_item_product_id	revenue
0	2013-07-25 00:00:00.0	24	319.96
1	2013-07-25 00:00:00.0	93	74.97
2	2013-07-25 00:00:00.0	134	100.00
3	2013-07-25 00:00:00.0	191	5099.49
4	2013-07-25 00:00:00.0	226	599.99
...
9115	2014-07-24 00:00:00.0	926	31.98
9116	2014-07-24 00:00:00.0	957	5399.64
9117	2014-07-24 00:00:00.0	1004	10399.48
9118	2014-07-24 00:00:00.0	1014	3148.74
9119	2014-07-24 00:00:00.0	1073	4199.79

9120 rows × 3 columns